


# CoR Training Pipeline



Key Hyperparameters:

$\lambda = 1.0$

$\mu = 0.5$

$v = 0.1$

$N = 8$  groups

$K = 3$  rounds

$\beta = 0.01$

$lr = 5e-7$

1,000 samples → 800x sample efficiency vs r1-distill